

Supervised Learning (Part I)

EN5422/EV4238 | Fall 2023

w02_supervised_1.pdf

(Week 2 – 1/2)

Contents

1	SUPERVISED LEARNING INTRO	2
1.1	SURVEY & QUIZZES.....	2
1.2	SUPERVISED LEARNING	2
2	EXAMPLE DATA.....	2
3	LINEAR MODELS	3
3.1	SIMPLE LINEAR REGRESSION.....	3
3.2	OLS LINEAR MODELS IN PYTHON	4
4	POLYNOMIAL INPUTS	7
4.1	ESTIMATION.....	7
5	K-NEAREST NEIGHBOR MODELS (K-NN)	12
5.1	EXAMPLE & NOTE	12
5.2	K-NN IN ACTION	13
6	PREDICTIVE MODEL COMPARISON (OR HOW TO CHOOSE THE BEST MODEL).....	14
6.1	PREDICTIVE MODEL EVALUATION	14
6.2	STATISTICAL DECISION THEORY	14
6.3	CHOOSE THE BEST <i>PREDICTIVE</i> MODEL	17

1 Supervised Learning Intro

1.1 Survey & Quizzes

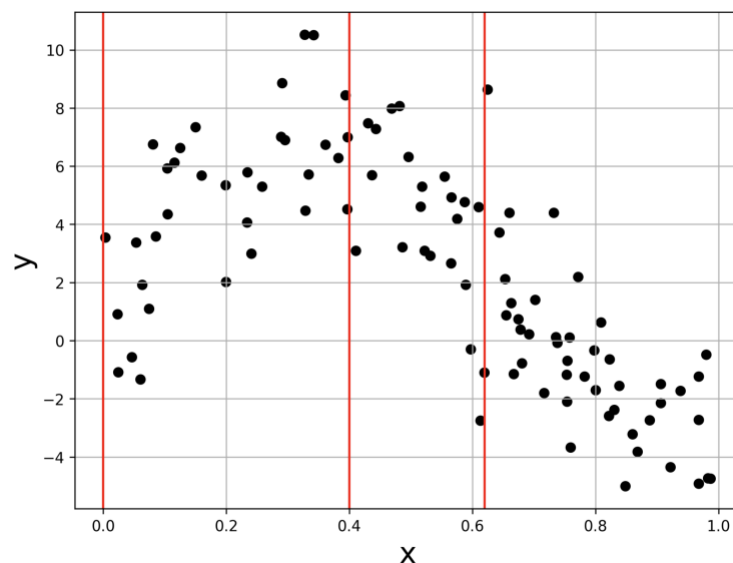
During the lecture, there will be a survey and quizzes. You can access the survey on [Slido.com](https://www.slido.com) using the # code provided during the lecture.

1.2 Supervised Learning

- In *supervised* learning, each observation can be partitioned into two sets: the predictor variables and the outcome variable(s).
 - Predictor variables are sometimes called 1) independent, 2) feature, and 3) predictor variables.
 - Outcome variables are sometimes called 1) target, 2) labels, 3) response, 4) dependent, and 5) reference variables.
- Usually, the predictor variables are represented by X and the response variables represented by Y .
- The goal in supervised learning is to find the patterns and relationships between the predictors, X , and the response, Y .
 - Usually, the goal is to predict the value of Y given X .
- Later in the course, we will explore the *unsupervised learning* topics of association analysis, density estimation, and clustering which do not have any outcomes (i.e., no Y 's)

2 Example Data

consider some data $D = \{(X_i, Y_i)\}_{i=1}^n$ with $Y_i \in \mathbb{R}$, $X_i \in [0,1]$ and $n = 100$.



Your Turn #1

The goal is to predict new Y values if we are given the X 's.

- If $x = 0.4$, predict Y .
- If $x = 0$, predict Y .
- If $x = 0.62$, predict Y .
- How should we build a model that will automatically predict Y for any given X ?

3 Linear Models

- Linear models refer to a class of models where the output (predicted value) is a linear combination (weighted sum) of the input variables:

$$f(x; \beta) = \beta_0 + \sum_{j=1}^p \beta_j x_j$$

where $x = [x_1, \dots, x_p]^T$ is a vector of features/variables/attribution and $\hat{Y}|x = f(x; \hat{\beta})$ is the predicted response at $X = x$.

- The coefficients (or weights), $\hat{\beta}$ are often selected by minimizing the squared residuals of the *training data* (may also be described as *ordinary least squares*; *OLS*)
 - But, there are other, and better, ways to estimate the parameters in linear regression that we will discuss later in the course. (e.g., Lasso, Ridge, Robust)

3.1 Simple Linear Regression

- single predictor variable $x \in \mathbb{R}$
- $f(x; \beta) = \beta_0 + \beta_1 x$
- Use *training data*: $D_{train} = \{(x_i, y_i)\}_{i=1}^n$
- OLS uses the weights/coefficients that minimize the residual sum of squares (RSS) loss function over the *training data*:

$$\hat{\beta} = \arg \min_{\beta} SEE(\beta)$$

where SSE is the sum of squared errors (also known as RSS):

$$SEE(\beta) = \sum_i^n (y_i - f(x_i; \beta))^2 = \sum_i^n (y_i - \beta_1 x_i)^2 = \sum_i^n \hat{\epsilon}_i^2$$

where, $\hat{\epsilon} = y_i - \hat{y}_i$ is the residual.

- The solutions are:

$$\hat{\beta}_0 = \bar{y} - \beta \bar{x}$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

- Definitions:

$$\begin{aligned} \text{MSE}(\beta) &= \frac{1}{n} \text{SSE}(\beta) \\ &= \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i; \beta))^2 \\ \text{RMSE} &= \sqrt{\text{MSE}} = \sqrt{\text{SSE}/n} \end{aligned}$$

3.2 OLS Linear Models in Python

3.2.1 Estimation with statsmodels

In **Python**, the library `statsmodels` fits an OLS linear model

```
import pandas as pd
import statsmodels.api as sm
import statsmodels.formula.api as smf

x = x.flatten()
y = y.flatten()

# Create a DataFrame
data_train = pd.DataFrame({'x': x, 'y': y})

# Fit the OLS model
model = smf.ols('y ~ x', data=data_train).fit()

# Print the summary
print(model.summary())

# Coefficients in DataFrame format
print(model.params)

# Other model properties
print("R-squared:", model.rsquared)
print("Adjusted R-squared:", model.rsquared_adj)
print("Standard error:", model.mse_resid ** 0.5)
print("F-statistic:", model.fvalue)
print("p-value of F-statistic:", model.f_pvalue)
print("AIC:", model.aic)
print("BIC:", model.bic)
print("Degrees of freedom of residuals:", model.df_resid)
print("Number of observations:", model.nobs)
```

OLS Regression Results						
Dep. Variable:	y	R-squared:	0.394			
Model:	OLS	Adj. R-squared:	0.388			
Method:	Least Squares	F-statistic:	63.68			
Date:	Sun, 20 Aug 2023	Prob (F-statistic):	2.80e-12			
Time:	08:20:29	Log-Likelihood:	-252.22			
No. Observations:	100	AIC:	508.4			
Df Residuals:	98	BIC:	513.6			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.8279	0.652	10.470	0.000	5.534	8.122
x	-8.6249	1.081	-7.980	0.000	-10.770	-6.480
Omnibus:	0.194		Durbin-Watson:		1.783	
Prob(Omnibus):	0.908		Jarque-Bera (JB):		0.051	
Skew:	-0.054		Prob(JB):		0.975	
Kurtosis:	3.023		Cond. No.		4.63	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Intercept 6.827862

x -8.624911

dtype: float64

R-squared: 0.3938577650885271

Adjusted R-squared: 0.3876726402424917

Standard error: 3.044531794318476

F-statistic: 63.678223947408824

p-value of F-statistic: 2.8006122101855e-12

AIC: 508.4368613513172

BIC: 513.6472017232934

Degrees of freedom of residuals: 98.0

Number of observations: 100.0

- statsmodels has an `ols` function, which includes the intercept by default.
 - Some examples of using formulas as well as getting the underlying X (model/design matrix) can be found [here](#).

3.2.2 Prediction with statsmodels

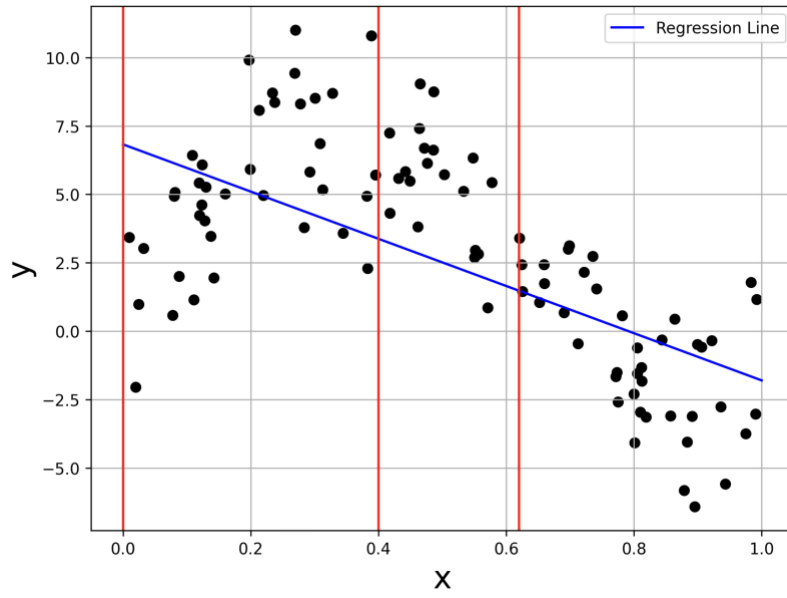
```
import numpy as np
import pandas as pd

# Add a constant term (for intercept) to x data
X = sm.add_constant(x)

# Fit the OLS model
model = sm.OLS(y, X)
result = model.fit()

# Create a prediction dataset
xseq = np.linspace(0, 1, 200) # sequence of 200 equally spaced values
                                # between 0 and 1
xeval = sm.add_constant(xseq)

# Predict using the fitted model
yhat1 = result.predict(xeval)
```



3.2.3 Questions

Your Turn #2

1. How did we do? If X_{new} is close to 0, or close to 0.4, or close to 0.62?
2. How to make it better?

4 Polynomial inputs

- In the *simple* linear regression model, we had 2 parameters that we needed to estimate, β_0 and β_1 . Thus, the **model complexity** is minimal.
 - The only thing simpler is an intercept only model.
- But the data appears to have a more *complex* structure than linear.
- A *parametric approach* to add complexity is to incorporate *polynomial terms* in the model.
 - A quadratic model is $f(x; \beta) = \beta_0 + \beta_1 x + \beta_2 x^2$

4.1 Estimation

- OLS uses the weights/coefficients that minimize the SSE loss function over the training data:

$$\begin{aligned}\hat{\beta} &= \arg \min_{\beta} SEE(\beta) \text{ Note: } \beta \text{ in this problem is a vector} \\ &= \arg \min_{\beta} \sum_{i=1}^n (y_i - f(x_i; \beta))^2 \\ &= \arg \min_{\beta} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i - \beta_2 x_i^2)^2\end{aligned}$$

4.1.1 Estimation

- Model**

$$f(x_i; \beta) = \mathbf{x}^T \beta$$

$$\mathbf{x} = \begin{bmatrix} 1 \\ x \\ x^2 \end{bmatrix} \quad \beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix}$$

Your Turn #3

Solve for $\hat{\beta}$ using matrix notation.

$$Y = \begin{bmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{bmatrix} \quad X = \begin{bmatrix} 1 & X_1 & X_1^2 \\ 1 & X_2 & X_2^2 \\ \vdots & \vdots & \vdots \\ 1 & X_n & X_n^2 \end{bmatrix} \quad \beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix}$$

Sol)

Loss function:

$$\begin{aligned}
 RSS(\beta) &= (Y - X\beta)^T(Y - X\beta) \\
 &= Y^T Y - \beta^T X^T Y - Y^T X \beta + \beta^T X^T X \beta \\
 &= Y^T Y - 2\beta^T X^T Y + \beta^T X^T X \beta \quad \text{Note: see below how } \frac{\partial \beta^T X^T X \beta}{\partial \beta} \text{ can be calculated.}
 \end{aligned}$$

$$\frac{\partial RSS(\beta)}{\partial \beta} = -2X^T Y + 2X^T X \beta = 0$$

$$\hat{\beta} = (X^T X)^{-1} X^T Y \quad \gg \text{ This is the least square solution.}$$

Given the expression $f(\beta) = \beta^T X^T X \beta$, where β is a column vector of dimension $p \times 1$ and X is a matrix of dimension $p \times 1$, you want to differentiate f with respect to β .

To differentiate this, we'll use the following properties of matrix derivatives:

1. $\frac{\partial}{\partial x} (X^T A X) = (A + A^T)X$ where A is a symmetric matrix
2. $\frac{\partial}{\partial x} (X^T A) = A$ where A is a matrix.

Let's differentiate the expression step by step:

Given:

$$f(\beta) = \beta^T X^T X \beta$$

Let $A = X^T X$. Then, $f(\beta) = \beta^T A \beta$.

Using the first property:

$$\frac{\partial f}{\partial x} = (A + A^T)\beta$$

Since $A = X^T X$ is symmetric (i.e., $A = A^T$), the derivative simplifies to:

$$\frac{\partial f}{\partial \beta} = 2A\beta$$

Substituting back for A :

$$\frac{\partial f}{\partial \beta} = 2X^T X \beta$$

So, the gradient of $f(\beta)$ with respect to β is $2X^T X \beta$.

4.1.2 Python implementation

In Python, the `OLS` function in `statsmodels.api` is a convenient way to get polynomial terms.

```
import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt

# Assuming sim_x and sim_y are already defined elsewhere, and you've already
# set up data_train as before
#-- Model Settings
n = 100
sd = 2
x = sim_x(n)
y = sim_y(x, sd)

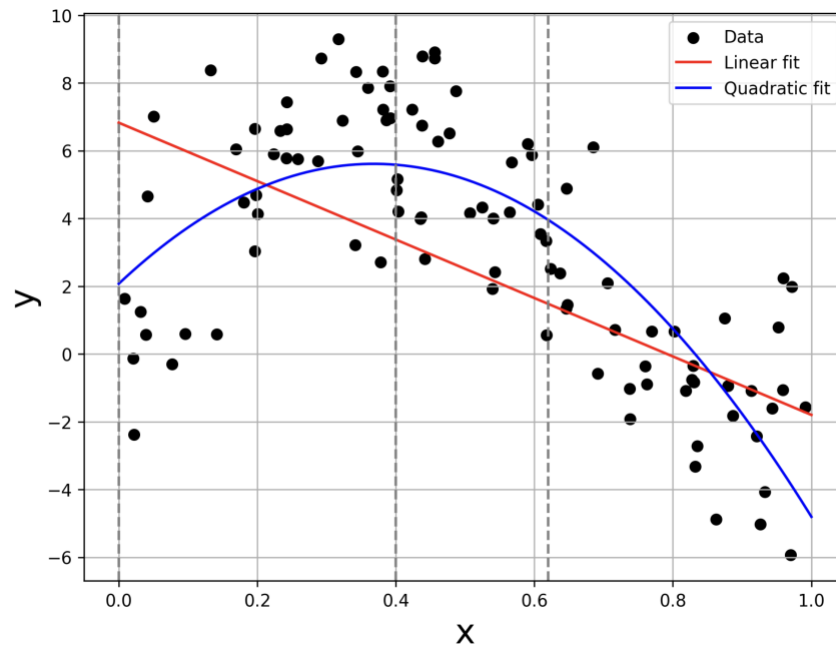
x = x.flatten()
y = y.flatten()

data_train = pd.DataFrame({'x': x, 'y': y})

# Create the polynomial features for the model
data_train['x2'] = data_train['x']**2
X = sm.add_constant(data_train[['x', 'x2']])

# Fit the quadratic regression model
model2 = sm.OLS(data_train['y'], X)
result2 = model2.fit()

# Create a prediction dataset
xseq = np.linspace(0, 1, 200) # sequence of 200 equally spaced values
# between 0 and 1
xeval2 = pd.DataFrame({'x': xseq, 'x2': xseq**2})
xeval2 = sm.add_constant(xeval2)
yhat2 = result2.predict(xeval2)
```

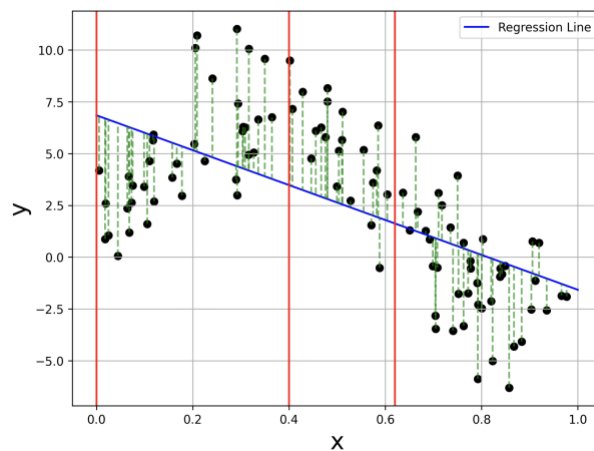


Your Turn #4


1. How did we do? If X_{new} is close to 0, or close to 0.4, or close to 0.62?
2. But does the quadratic model fit better *overall*?
3. What is the *complexity* of the quadratic model?
(The effective degrees of freedom is essentially the number of parameters you're estimating.)

Degree	MSE	# of params
1	8.3	2
2	5.5	3

Linear Fit



Quadratic Fit



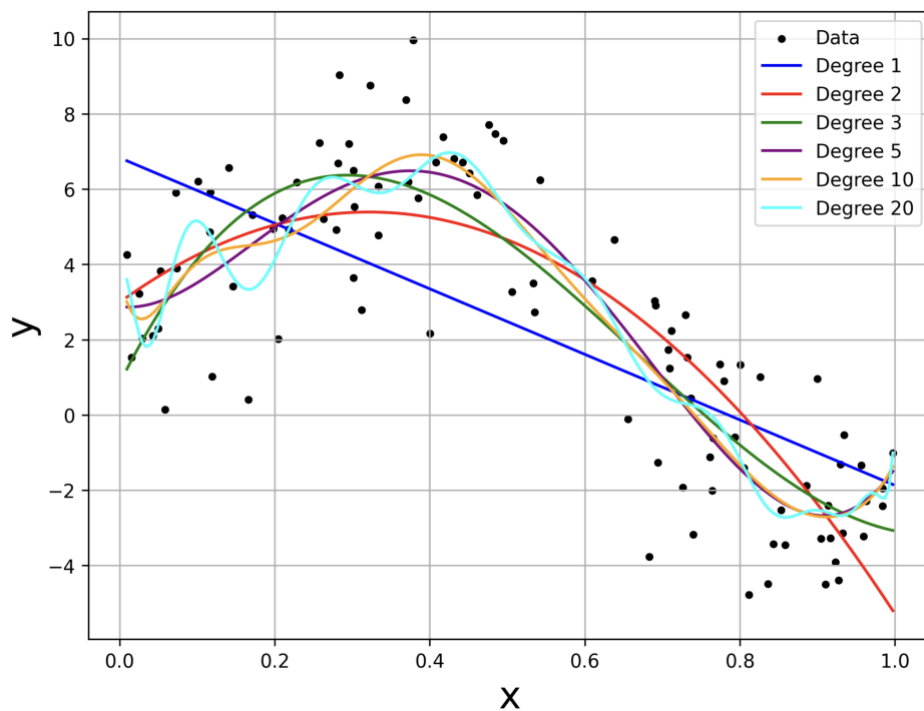
As kids always reason, “if a little is good, than a lot much be better”> So why not try more complex models by increasing the polynomial degree.

- Polynomial of degree d

$$f_{poly}(x; \beta, d) = \beta_0 + \sum_{j=1}^d \beta_j x^j$$

Degree	MSE	# of params
1	7.5	2
2	4.5	3
3	3.7	4
5	3.3	6
10	3.2	11
20	3.0	21

And its always good to observe the plot.



- For degree=20, the behavior at the end points are a bit erratic.
- Using a higher degree would further reduce the MSE, but the fitted curve would be more “complex” and may not be a s good for new data.

5 k -nearest neighbor models (k -NN)

- The k -NN method is a non-parametric *local* method, meaning that to make a prediction $\hat{y}|x$, it only uses the training data in the vicinity of x .
 - contrast with OLS linear regression, which uses all x 's to get prediction.
- The model is simple to describe:

$$f_{\text{knn}}(x; k) = \frac{1}{k} \sum_{i: x_i \in N_k(x)} y_i$$

$$= \text{Avg}(y_i \mid x_i \in N_k(x))$$

- $N_k(x)$ are the set of k nearest neighbors to x
- only the k closest y 's are used to generate a prediction
- it is a *simple mean* of the k nearest observations

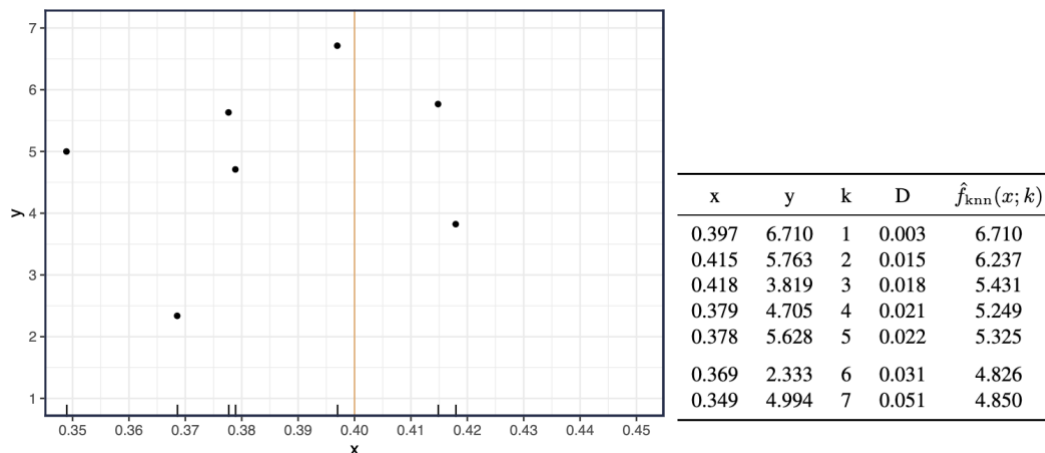
Your Turn #5

What is the estimate $f_{\text{knn}}(x; k = n)$?

5.1 Example & Note

5.1.1 Example

Let's zoom in on the region around $x=0.4$



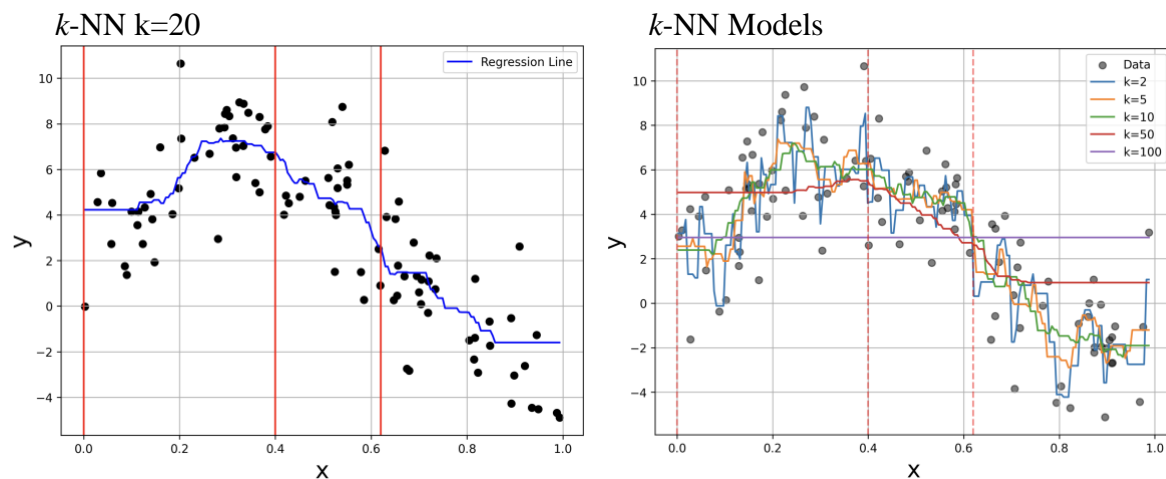
5.1.2 Note about k -NN

- A suitable *distance* measure (e.g., Euclidean) must be chosen.
 - And predictors are often *scaled* (same standard deviation or range) so one variable doesn't dominate the distance calculation.

- Because the distance to neighbors grows exponentially with increased dimensionality/features, the *curse of dimensionality* is often referenced with respect to k -NN.
 - This means that in high dimensions most *neighbors* are not very close and the method becomes less *local*.
- One computational drawback of k -NN methods is that all the training data must be stored in order to make predictions.
 - For large training data, may need to sample (or use prototypes)
- The complexity of a k -NN model increases as k decreases.
- The least complex model, which is a constant, occurs when $k = n$.
- The most complex model is when $k = 1$.
- The **effective degrees of freedom** or *edf* for k -NN model is n/k .
 - this is a measure of the model complexity. It is approximately the number of parameters that are estimated in the model (to allow comparison with parametric models)

5.2 k -NN in action

In Python, the function `KNeighborsRegressor` from `sklearn.neighbors` library will fit a k -NN regression model. Here is a $k = 20$ nearest neighbor model.



5.2.1 k -NN in action

k	MSE	edf
100	7.5	1
50	4.5	2
10	3.7	10
5	3.3	20
2	3.2	50

6 Predictive Model Comparison (or how to choose the best model)

6.1 Predictive Model Evaluation

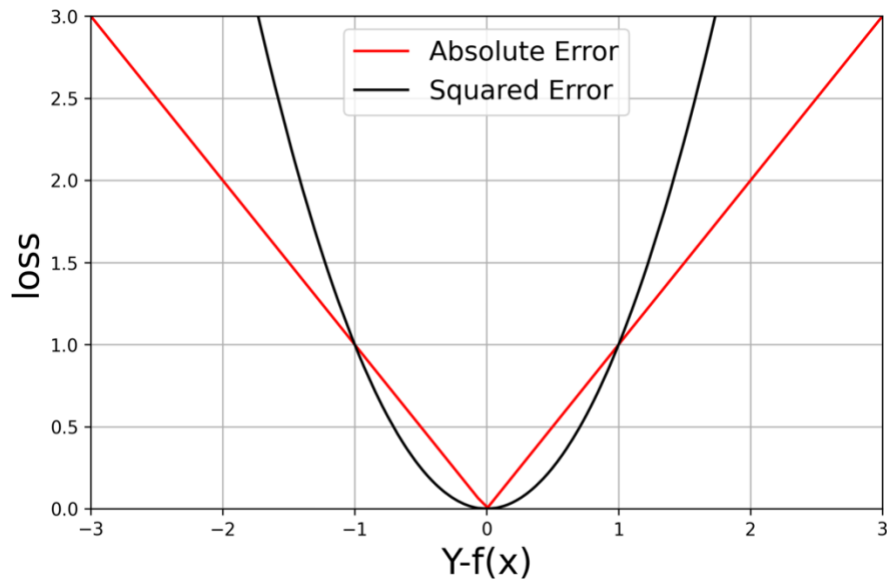
Our goal is prediction, so we should evaluate the models on their *predictive performance*.

- We need to use hold-out data (i.e., data not used to fit the model) to evaluate how well our models do in prediction.
- Call these data test data $D_{\text{test}} = \{(X_j, Y_j)\}_{j=1}^J$
 - Note: assume that the test data comes from the same distribution as the training data
 - Or $P_{\text{test}}(X, Y) = P_{\text{train}}(X, Y)$
 - **both** Y and X from same distribution
- Late in the course we will cover ways to do this when we only have training data (e.g., cross-validation)
- but for today, we have an unlimited amount of *test data* at our disposal (since we know how the data were generated)

6.2 Statistical Decision Theory

- In a prediction context, we want a *point estimate* for the value of an unobserved random variable (r.v.) $Y \in \mathbb{R}$ given an input feature $X \in \mathbb{R}$.
- Let $f(X)$ be the prediction of Y given X .
- Define a *loss function* $L(Y, f(X))$ that indicates how bad it is if we estimate the value Y by $f(X)$.
 - e.g., Y is the number of rainfall events at GIST and X is the day of week from August to October.
 - If we guess $f(X) = 2$, but there are really $Y = 1$, how bad would that be?
- A common loss function is *squared error*:

$$L(Y, f(X)) = (Y - f(X))^2$$



- The best model is the one that minimizes the *expected loss* or **Risk** or **Expected Prediction Error (EPE)**

$$Risk = EPE = E[loss]$$

- For *squared error*, the *risk* for using the model f is:

$$R(f) = E_{XY}[L(Y, f(X))] = E_{XY}[(Y - f(X))^2]$$

where the expectation is taken with respect (w.r.t.) to the *test values* of X and Y .

- Note under squared error loss, the risk is also known as the *mean squared error* (MSE)
- To simplify a bit, let's examine the risk of model f at a given fixed input $X = x$. This removes the uncertainty in X , so we only have uncertainty coming from Y .

$$\begin{aligned} R_x(f) &= E[L(Y, f(x)) | X = x] \\ &= E[(Y - f(x))^2 | X = x] \end{aligned} \quad \text{for squared error loss}$$

where the expectation is taken with respect to $Y | X = x$.

- The best prediction $f^*(x)$, given $X = x$, is the value that minimizes the risk:

$$\begin{aligned} f^*(x) &= \arg \min_c R_x(c) \\ &= \arg \min_c E[(Y - c)^2 | X = x] \end{aligned}$$

Your Turn #6

What is the optimal prediction at $X = x$ under the squared error loss?

- i.e., find $f^*(x)$.

6.2.1 Squared Error Loss Functions

- Conclusion: If quality of prediction is measured by squared error, then the best predictor is the (conditional) expected value $f^*(x) = E[Y|X = x]$.
 - And the minimum Risk/MSE is $R_x(f^*) = V[Y|X = x]$.
- **Summary:** Under *squared error loss* the Risk (at input x) is

$$\begin{aligned}
 R_x(f) &= E_Y[L(Y, f(X))|X = x] \\
 &= E_Y[(Y - f(x))^2|X = x] \quad \text{using squared error loss} \\
 &= \text{Var}[Y|X = x] + (E_Y[Y|X = x] - f(x))^2 \\
 &= \text{Irreducible Variance} + \text{model squared error}
 \end{aligned}$$

$$\text{Note: } \text{Var}[Y|X = x] = E_Y[Y^2|X = x] - (E_Y[Y|X = x])^2$$

6.2.2 k -NN and Polynomial Regression

- The k -NN model estimates the conditional expectation by using the data in a *local region* around x :

$$\hat{f}_{knn}(x; k) = \text{Ave}(y_i | x_i \in N_k(x))$$

This assumes that the true $f(x)$ can be well approximated by a *locally constant* function.

- Polynomial (linear) regression, on the other hand, assumes that the true $f(x)$ is well approximated by a *globally polynomial* function.

$$\hat{f}_{poly}(x; d) = \beta_0 + \sum_{j=1}^d \beta_j x^j$$

6.2.3 Empirical Risk

- The actual Risk/EPE is based on the expected error from *test data* (out-of-sample), or data that was not used to estimate \hat{f} .

$$\begin{aligned} R(f) &= E_{XY}[L(Y, f(X))] \\ &= E_{XY}[(Y - f(X))^2] \quad \text{for squared error loss} \end{aligned}$$

where X, Y are from $\Pr(X, Y)$ (i.e., test data).

- But is it a bad idea to choose the best model according to empirical risk or training error?

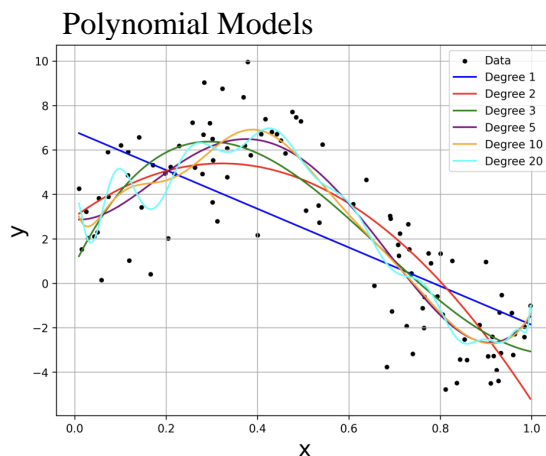
$$\begin{aligned} \hat{R}_n(f) &= \frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i)) \\ &= \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2 \quad \text{for squared error loss} \end{aligned}$$

6.3 Choose the best *predictive* model

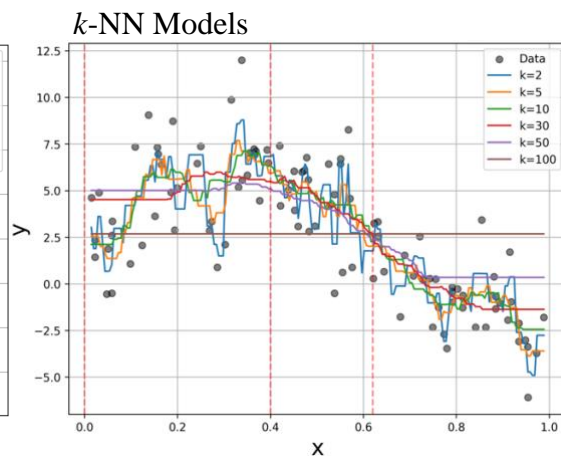
Your Turn #7

Which model will you choose?

Enter your answer on *slido*.



Degree	MSE	# of params
1	7.5	2
2	4.5	3
3	3.7	4
5	3.3	6
10	3.2	11
20	3.0	21



k	MSE	edf
100	12.7	1
50	6.6	2
30	4.6	3.3
10	3.9	10
5	3.7	20
2	2.2	50